

Problem Set #2 Solutions

Out: May 22

Due: June 1

1. (a) We can show that the inequality in the question holds by considering the structure of M_1 and M_2 . The breakdown of matrix of M into M_1 and M_2 looks like

$$\begin{pmatrix} \times & X_1 \times T & \times \\ S \times Y_1 & S \times T & S \times Y_2 \\ \times & X_2 \times T & \times \end{pmatrix}$$

where the highlighted row is matrix M_1 , the highlighted column is the matrix M_2 and X_1, S, X_2 with Y_1, T, Y_2 are partitions of X and Y respectively. Note that due to the structure of given matrices, the standard rank of M_1 is bound above by $|S|$ and the standard rank of M_2 is bound above by $|T|$, and the same bounds apply to triangle-rank.

Note also that $S \times T$ is a zero matrix - this means that regardless of how rows or columns from $S \times T$ are permuted within M_1 or M_2 , they will not change the triangle-rank. As such, there is no reason to permute rows indexed by S in M_2 and no reason to permute columns indexed by T in M_1 . This creates a convenient “mask” that lets us permute rows and columns of M and build desired square submatrices inside M_1 and M_2 without undoing any progress.

Now we can put the two points from above together to define a procedure that proves the desired inequality. We start off with M . Next we permute rows and columns of M to build the largest square submatrix with ones in diagonals on zeros below the diagonal in both M_1 and M_2 (showing triangle-rank of these matrices). As discussed previously, we can do that without any clashes.

Denote such square submatrices S_1 in M_1 and S_2 in M_2 . Finally, we permute columns of resultant M to put S_1 into the top left corner of block $S \times Y_2$, and then permute rows of resultant M to put S_2 into the bottom right corner of block $X_1 \times T$. As a result of this operation, we have transformed M in such a way that it now contains a square matrix that is a tensor product of S_1 and S_2 (with the exception of the top right corner which contains arbitrary values and doesn't affect triangle-rank). Hence we have shown that if matrices M_1 and M_2 have triangle-rank of m and n respectively, we can permute rows and columns of M to build a square submatrix of side length $m + n$ with ones on the diagonal and zeroes below the diagonal, showing that the triangle-rank of matrix M must be at least $m + n$.

- (b) We can prove the desired bound using induction on triangle-rank. The high level overview of the protocol is as follows. Before communication begins, Alice and Bob fix the smallest non-disjoint 0-cover of M_f . Denote the set of all rectangles in that cover as C_0 . Note that $\log_2 |C_0| = N^0(M_f)$, so it takes $N^0(M_f)$ bits to communicate the index of each rectangle in C_0 . Alice and Bob maintain sets A and B respectively, where A is the set of rows that are still in the search space for Alice and B is the set of columns that are still in the search space for Bob. Both A and B start off as the full matrix M_f . On input x^* for Alice, which corresponds to row x in the matrix, and input y^* for Bob, which corresponds to column y of the matrix, Alice and Bob proceed as follows.

Alice begins the communication. She iterates over all $S \times T$ rectangles $R_A \in C_0$ that satisfy the three conditions below until she runs out of suitable rectangles. Conditions for each R_A are:

- i. R_A intersects her row x .
- ii. At least one of the rows of R_A overlaps with the set A .
- iii. Denote M^A as the matrix representing Alice's search space, $A \times Y$. Restricting Alice's search space A to the rows that only appear in both A and S (intersection) should cut the rank of M_A by at least a factor of 2.

When Alice finds such a rectangle, she sends its index over to Bob. Bob replies with 0 if his column y also intersects this 0-rectangle or 1 otherwise. Note that if Bob replies 0, protocol can terminate since Alice and Bob found a 0-rectangle that they both intersect, so $f(x, y) = 0$. Otherwise, if Bob replies 1, Alice restricts her A to the rows that appear in intersection of A and S and continues the search. Note that since Bob also knows the index of the rectangle, he know what Alices A looks like.

Because of the conditions we imposed on suitable R_{AS} , Alice cuts the triangle-rank of $A \times Y$ by at least a factor of 2 on every iteration. Additionally, every iteration takes $N^0(M_f) + 1$ bits of communication. Once Alice runs out of suitable rectangles, she signals Bob to take over. Bob now begins the same iteration process, but considering columns instead of rows and using $A \times B$ as his seach space, cutting triangle-rank at every step by shrinking B . If at some point Alice confirms that she's in the 0-rectangle that Bob specified, the protocol can terminate since $f(x, y) = 0$. Otherwise, Bob runs out of suitable rectangles and signals this fact to Alice, so both parties can conclude that $f(x, y) = 1$.

The correctness of this protocol is easy to show. If Alice transmits an index of a 0-rectangle and Bob accepts it (or vice versa), the conclusion is trivial. When Alice runs out of rectangles and has to signal Bob, there are 2 cases: $f(x, y) = 1$ and $f(x, y) = 0$. If $f(x, y) = 1$, then regardless of which rectangle Bob picks, Alice will never accept it, and Bob will eventually run out of rectangles so parties will conclude $f(x, y) = 1$. On the other hand, if $f(x, y) = 0$, then there must exist some $S \times T$ 0-rectangle that both x and y intersect. We know that Alice was not able to find this rectangle, so it must be the case that $\text{triangle-rank}((A \cap S) \times Y) \geq \frac{1}{2} \text{triangle-rank}(A \times Y)$. By rearranging the inequality from part (a), we can show that $\text{triangle-rank}(A \times T) \leq \frac{1}{2} \text{triangle-rank}(A \times Y)$, which means that Bob must consider this rectangle during his search. Since Bob is going to transmit its index, Alice must accept it and both parties will conclude $f(x, y) = 0$.

We can now use this protocol outline to prove the desired bound on communication complexity using induction on triangle-rank.

- **Inductive hypothesis:** Assume that the inequality

$$D(M_f) \leq (\log_2(\text{triangle-rank}(M_f) + 1) + 1)(N^0(M_f) + 1)$$

holds for all $M_f = A' \times B'$ such that $\text{triangle-rank}(A' \times B') < t$.

- **Base case:** $\text{triangle-rank}(M_f) = 0$. In this case, we have an all 0s matrix, so there is a 0-cover of size 1, hence $N^0(M_f) = \log_2(1) = 0$. Using our protocol, Alice transmits $N^0(M_f)$ bits, or, in other words, doesn't transmit anything since the index is trivial. Bob replies with 0 to acknowledge that he's in the same rectangle. In total, we have 1 bit of communication, and

$D(M_f) = 1$. Indeed, the formula holds for our base case:

$$\begin{aligned}
 D(M_f) &\leq (\log_2(\text{triangle-rank}(M_f) + 1) + 1)(N^0(M_f) + 1) \\
 &= (\log_2(0 + 1) + 1)(0 + 1) \\
 &= (0 + 1)(1) \\
 &= 1
 \end{aligned}$$

$\log_2(\dots + 1)$ guarantees that we don't hit $\log_2 0$.

- **Inductive step:** Need to show that formula holds for $M_f = A \times B$ s.t. $\text{triangle-rank}(A \times B) = t$. As per the protocol we outlined above, Alice starts searching for a suitable 0-rectangle from the cover first.

Case 1: Alice finds some $S \times T$ rectangle R_A and sends its index over to Bob. If Bob replies with 0, parties have found a common 0-rectangle and the protocol can terminate with output 0 after having exchanged only $N^0(M_f) + 1$ bits. Clearly, the inequality holds since we under-shoot our budget by having only 1 round. If Bob replies with 1, Alice has to continue the search. At this stage, we have expended $N^0(M_f) + 1$ bits, but we have also at least halved the the triangle-rank of our search space, which means that we're still within our budget of $\log_2(\text{triangle-rank}(A \times B) + 1) + 1$ rounds. We decrease the triangle rank of our search space, and we know that the inequality holds for values $< t$, so it must also hold for t . Note that we're adding 1 to the result of $\log_2(\dots)$ - this equivalent to "rounding up" our upper bound in the inequality, which is necessary since the actual number of bits exchanged in the protocol is discrete.

Case 2: Alice doesn't find a suitable rectangle, so she has to tell Bob that he needs to begin the search. She hands the control over to Bob. Earlier, we have proved the correctness of the algorithm, so Bob should be able to carry on. What we have now is logic from Case 1 carried out by Bob instead of Alice. Clearly, the inequality still holds.

- By induction, the proposed inequality holds for all values of $t \geq 0$.
2. We can prove that the log-rank conjecture holds for our notion of rank by reducing the problem of calculating the output $f(x, y)$ to an instance of CIS_G . I will assume that matrix multiplication is done using the standard dot product of rows and columns (as opposed to $\langle \cdot, \cdot \rangle \pmod 2$).

Denote the Boolean-rank of an $N \times N$ matrix M as r . Then there must exist an $N \times r$ Boolean matrix U and an $r \times N$ Boolean matrix V such that $M = UV$. Note that if we use standard matrix multiplication, this is only possible when $\sum_{q=0}^r U_{iq} * V_{qj} \in \{0, 1\}$ for all $i, j \in \{1, \dots, N\}$. This, in turn, implies that every row x in U and every row y in V share at most one index p such that $x_p * y_p = 1$ (for all remaining indices p' , it must be the case that $x_{p'} * y_{p'} = 0$). We can exploit this fact in our reduction to CIS_G .

Given a function $f : X \times Y \rightarrow \{0, 1\}$ and the corresponding M_f with $\text{Boolean-rank}(M_f) = r$, we can construct a graph G as follows. G will have r nodes, denoted a_i for $i \in \{1, \dots, r\}$. The node a_i corresponds to i th column of U and i th row of V . We then add edges as follows. For every row x in U , we treat x as a bitmap of nodes from G , where $x_i = 1$ means that the bitmap includes node a_i . We connect all nodes a_i that appear in the bitmap x together, so that they form a clique. Graph G is now ready to be used in CIS_G .

The protocol works as follows. Before communication begins, both Alice and Bob know G , U and

V . On input x^* that corresponds to row x in M_f , Alice looks at row x in U and treats it as a bitmap of nodes from G that she has to choose as her clique, C . By definition of G , her choice of C is guaranteed to be a clique.

On input y^* that corresponds to column y in M_f , Bob looks at column y of V and treats it as a bitmap of nodes from G that he has to choose as his independent set, I . We can show by contradiction that Bob's choice of I must be an independent set. Assume that I is not an independent set. This means that there is an edge between some nodes in I . By construction of G , this in turn implies that matrix U has some row y^+ that shares more than one index i with x such that $y_i^+ * x_i = 1$, which implies that not all matrices in relationship $M = UV$ are Boolean matrices, which is a contradiction. Hence I must be an independent set.

Note that no communication has yet occurred. Alice and Bob now solve problem CIS_G using inputs C and I , using output of $CIS_G(C, I)$ as the result of running $f(x^*, y^*)$. Proof of correctness is trivial: if CIS_G outputs 1, then C and I share some common node a_i . This implies that $x_i * y_i = 1$, $\langle U_x, V^y \rangle = 1$. The last expression is equivalent to calculating the value of the element in M_f on which row x and column y intersect each other, i.e. the output of $f(x, y)$. Similarly, if CIS_G outputs 0, we have $\langle U_x, V^y \rangle = 0$ and $f(x, y) = 0$.

We have shown in class that communication complexity of CIS_G is bounded above by $O(\log^2 n)$, where n is the size of the graph. In our case, size of the graph is $r = \text{Boolean-rank}(M_f)$, so we can say that

$$D(M_f) \leq O(\log^2 \text{Boolean-rank}(M_f))$$

This proves that log-rank conjecture holds for the proposed notion of rank.

3. (a) First, we can observe that the definition of discrepancy implies that we can invert Boolean outputs of the function without changing the discrepancy, since we're using the absolute difference between probabilities.

We can start by looking at the full $X_1 \times \dots \times X_k$ cylinder first. With this choice of S , the calculation for

$$\left| \Pr_{\mu}[f(x_1, \dots, x_k) = 0 \wedge (x_1, \dots, x_k) \in S] - \Pr_{\mu}[f(x_1, \dots, x_k) = 1 \wedge (x_1, \dots, x_k) \in S] \right|$$

reduces into just

$$\left| \Pr_{\mu}[f(x_1, \dots, x_k) = 0] - \Pr_{\mu}[f(x_1, \dots, x_k) = 1] \right|.$$

It's clear that the maximum possible value is 1, when the function either outputs always 1 or always 0. When we don't have this trivial case, the two terms begin to cancel each other out. That is, if we WLOG assume $\Pr_{\mu}[f(\dots) = 1]$ is the largest of the two, we know that

$$\left| \Pr_{\mu}[f(\dots) = 1] \right| > \left| \Pr_{\mu}[f(\dots) = 0] - \Pr_{\mu}[f(\dots) = 1] \right|.$$

Due to this fact, we can aim to maximise the value by somehow getting rid of the smaller component. We can rewrite the definition of discrepancy as follows:

$$\text{disc}_{\mu}(f) = \max_S \left| \Pr_{\mu}(f(x^k) = 1) \Pr_{\mu}(x^k \in S \mid f(x^k) = 1) - \Pr_{\mu}(f(x^k) = 0) \Pr_{\mu}(x^k \in S \mid f(x^k) = 0) \right|$$

From now on, we will WLOG focus on 1-monochromatic cylinder intersections and assume given f has nondeterministic complexity t . Identical argument applies to the co-nondeterministic version. If we WLOG consider only 1-monochromatic cylinder intersections, it's clear the second term becomes zero and we're left off with just $\Pr_\mu(f(x^k) = 1) \Pr_\mu(x^k \in S \mid f(x^k) = 1)$.

We're given that f has nondeterministic complexity t , which implies that there exists a non-disjoint cover of f with 1-monochromatic cylinder intersections, with the size of this cover being 2^t . Note that because this cover can be non-disjoint, total sum of probabilities of an element belonging to these cylinders can be greater than 1. Since these cylinder intersections cover *all* 1s and there are 2^t of them, they must include at least one cylinder intersection S such that $\Pr_\mu(x^k \in S) \geq 2^{-t}$. Otherwise, we would have $\Pr_\mu(x^k \in S) < 2^{-t}$ for all of the 2^t 1-cylinder intersections in the cover, which is a contradiction since the sum of probabilities will be less than 1. As a result, we can conclude that there exists a 1-monochromatic cylinder intersection S such that $\Pr_\mu(x^k \in S \mid f(x^k) = 1) \geq 2^{-t}$.

The function f and given distribution μ are fixed, so we can conclude that

$$\Omega(\Pr_\mu(f(x^k) = 1) \Pr_\mu(x^k \in S \mid f(x^k) = 1))$$

simplifies to just $\Omega(\Pr_\mu(x^k \in S \mid f(x^k) = 1))$, and, finally, using the result from above, to $\Omega(2^{-t})$ when we use the largest 1-monochromatic cylinder intersection. This shows that the discrepancy is bounded below by $\Omega(2^{-t})$ given that nondeterministic or co-nondeterministic communication complexity of f is t .

- (b) We can view the input for the problem DISJ_k as a $k \times n$ Boolean matrix, where each player i holds the string corresponding to i th row. In this formulation, the co- DISJ_k problem reduces to finding a column of all 1's (which implies a global intersection).

An all powerful prover can point out the index of the column that contains all 1's, which would require $\log n$ bits to represent. Then, the first player can check everyone else's values in that column and output 1 if they are indeed all 1's. The second player can check first player's value and output 1 in the same case, requiring $O(\log n)$ bits in total to solve co- DISJ_k .